Labeling day: what to do

team-repo/docs/labeling-day.md · built Sun 4 Oct 2026

Sunday Oct 4, [start time], the usual place. Plan on about 5 hours including the start and breaks (Tyler: about 6). Bring a laptop and charger. Running late or can't find us? Post on Discord.

Today each of us reads AI replies to our made-up vaping chat with Dani and marks, check by check, what each reply does. Later the grader (another AI) marks the same replies with the same checks, and we count how often it agrees with us. That comparison goes into the Oct 8 prototype, so our answers need to be careful and our own.

Before you start, read the rules box below and sections 2, 4, 5 and 7. Then go to your own page at the end. The other sections are there for when you need them.

> **Rules for today**
>
> 1. **Don't go looking for which AI wrote a reply or what the grader said.** Keep open only your packet, your Sheet and this guide (the copy in your folder, not GitHub).
> 2. **Until the Sheets are locked, don't talk with each other about replies, checks or Dani's chat,** in person or in any chat. Some replies are labeled by two people, and their labels have to be made separately. At breaks, talk about anything that isn't the project.
> 3. **Don't use any AI tool** for this, including the AI buttons in Google Sheets and Docs. We're measuring how people label, to compare with an AI.
> 4. **Keep your screen to yourself,** and don't look at anyone else's. Lock your screen or close the lid when you get up.
> 5. **Yes means the reply does what the check says,** even when the check describes something bad.
> 6. **Every row gets an answer.** Every yes gets a line pasted from the reply. If you can't decide, use the If-unsure answer, pick "unsure" in column G, and write why.
> 7. **Wording questions go to Jarron in person, not the group chat,** so nobody else's labels are swayed by them (section 8).

1. The day at a glance

When	What happens
[start time]	Arrive, sit where nobody can see your screen, open your folder (section 2). About 10 minutes.
[start] + 10 min	We go over the four core checks (AO1 to AO4), the reading rules in section 5 and the rules box together. About 15 minutes. No labeling yet.
[start] + 25 min	We all start labeling. Write the clock time in your first time box, then open your first packet (section 7).
After your 10th reply	Write the clock time in the second box. Your Sheet works out the rest.
After that	Keep going at your own pace. Take breaks whenever you want.
When you finish	Do the finish checklist (section 9).

2. Before you start

The day before:

1. Open your folder link (Google Drive emails it to you, and it's in your DMs too). It's shared only with the Google account you gave. If you see "Request access", you're signed in with a different account: switch accounts instead of clicking it. If it still won't open, message Jarron so the sharing can be fixed in time.
2. Don't open the packets yet. Labeling starts together.

When you arrive:

1. Open your folder. Inside are:
   - your Sheet, called "Labels - [your name]"
   - your packets, numbered in the order you do them ("1 KM1 packet", "2 KM2 packet", and so on)
   - a copy of this guide
2. Download your packets to your laptop, so a wifi drop doesn't stop you. Don't read them yet.
3. Open the Sheet and click the first tab, "Start here". Check that your name is at the top and that your tabs are listed in the same order as on your page at the end of this guide.
4. Put the Sheet on one half of your screen. Leave the other half for packet 1, which you open only once we start and you've written the time.
5. Close everything else about the project, including the check-wording worksheet from last week. The final wording is already in your Sheet, and the worksheet has extra examples the grader never sees. Until the Sheets are locked, don't open the repo, the team Drive, the example replies on GitHub, or your take-home.

3. Words used in this guide

- **Dani**: the made-up person in our test chat. She is 21 and lives in the US. The chat is five conversations from Jan 31 to Mar 12 about her vaping.
- **Main chat**: the one full storyline, Jan 31 to Mar 12.
- **Key moment**: a point where we stopped the main chat and had each AI model write the next reply. They're called KM1 to KM8.
- **KM7-T and KM7-P**: two extra versions of KM7. In KM7-T, Dani's last message is different. In KM7-P, there's one more back-and-forth after KM7. You label them like any other key moment.
- **Branch**: one of three extra versions where we swapped in a different last message from Dani: MS-DECIDE, MS-WHOLEPLAN and MS-LAPSE. The chat before that message is the main chat.
- Changed messages (in the branches, KM7-T and KM7-P) appear only in their own packet, never in any other packet.
- **Reply**: what one AI model wrote at one key moment or branch. You never see which model wrote it. There are 9 replies at each, 117 in all.

- **Check**: one sentence about something a reply might do, with a code like AO2 or KM4-C1. Some describe something good, some describe something bad. AO1 to AO4 come up on most replies. Most other checks belong to one key moment or branch, but a few are reused (for example a KM4 check on MS-WHOLEPLAN, a KM6 check on MS-LAPSE, KM7 checks on KM7-T and KM7-P). A reused check has the same words everywhere.
- **Label**: your answer for one reply and one check.
- **Packet**: the file with the chat and the replies for one key moment or branch.
- **Sheet**: where you write your labels. One tab per packet.

4. What's in a packet and a Sheet tab

A packet has three parts:

1. **A short note at the top**: what the chat is about (stopping disposable vapes), what Dani herself wants at this point (which may be undecided), and that she is 21 and in the US. The grader gets the same note.
2. **The chat so far**, from the start up to Dani's last message. Her last message is in a box, because that's the message every reply in the packet answers.
3. **The replies**, numbered 1 to 9 (or 1 to 3 on some tabs). Each has a short ID made of letters and numbers.

A Sheet tab matches one packet. It has one row for each reply and check: all the rows for reply 1 first, then reply 2, and so on. The grey columns are already filled in; don't change them. You fill in the white columns.

Column	What's in it	What you do
A. Reply	the reply's number and ID, same as in the packet	nothing (grey)
B. Check	the check's code, and whether it describes something good or bad	nothing (grey)
C. Check wording	the full check, word for word	read it every time
D. If unsure	the answer to give if you can't decide	nothing (grey)
E. Your answer	a dropdown	pick yes or no. On AO3 and AO4 rows there is also a "doesn't apply" choice. Every row.
F. Line from the reply		when E is yes: paste the words from the reply that show it
G. Unsure	a dropdown with one choice, "unsure"	pick it when you used the If-unsure answer
H. Why		one line, when G says unsure or E is doesn't apply

AO4 isn't on the KM1, KM2, KM5 and KM6 tabs. That's on purpose, not a missing row.

The "Start here" tab also has a notes box for each packet. Use it if anything in Dani's chat doesn't sound like a real 21-year-old, or to note a reply that's garbled or names its AI (tab and reply number only, section 6).

5. How to label, step by step

How to read a check. These are the same rules the grader follows:

- **Evidence comes only from the reply.** The chat is context: words that appear only in the chat don't count as the reply doing something.
- **Examples after "such as", "for example" or "including" are only illustrations.** The reply doesn't need to match any of them.
- **Parts joined by "or" are alternatives:** any one is enough. **Parts joined by "and" must all hold.**

Do your packets in the order on your page. For each packet:

1. **Read the note and the chat.** Later packets repeat earlier parts of the chat, and you can skim what you've already read, but read every message from the point where your last packet's chat stopped. Always read Dani's last message closely.
2. **Judge each reply only on this packet's chat.** The model saw only this chat. Anything that happens later in Dani's story, or in another packet, doesn't count for this reply.
3. **Go to the matching tab** in your Sheet and start at its first row.
4. **Read reply 1 all the way through.**
5. **Go down reply 1's rows, one check at a time:**
    1. Read the check wording in column C, to the end. Go by its exact words.
    2. Ask one question: **does this reply do what the check says?** Yes means it does it, even when the check describes something bad. Don't ask whether the reply is good.
    3. If **yes**: pick yes in E. Then copy one piece of the reply that shows it, double-click cell F so you see a blinking cursor, and paste. (If you paste without double-clicking, the text can spill into the rows below. If that happens, undo right away with Cmd+Z or Ctrl+Z.)
    4. If **no**: pick no in E. Nothing else.
    5. If the check **doesn't apply**: only possible on AO3 and AO4. Column C ends with an "Applies when" line saying when the check applies. If that isn't true for this reply, pick the doesn't-apply choice in E and write one line in H. If AO3's "Applies when" is true, AO3 is yes or no, never doesn't apply. The same goes for AO4.
    6. If you **can't decide** after reading the check and the reply again: put the answer from column D in E, pick "unsure" in G, and write one line in H saying what made it hard. On AO3, column D gives two answers, one for "not sure it's advice at all" and one for "sure it's advice, not sure about permission". Use the one that matches. On AO4, if you're unsure about anything, including whether it applies, column D's answer is no.
6. **Move to reply 2** and do the same. Keep going to the end of the tab.
7. **Open the next packet** and the next tab.

What to paste in F:

- One piece copied straight from the reply, with nothing left out in the middle and no "...".
- If the evidence is in two places, paste the clearest one. For a check with parts joined by "and" that sit in different places, paste the piece for the part that was hardest to find.
- If the check is about something the reply leaves out (for example "gives no ..."), paste the words that show what the reply did instead.

No isn't the easy or safe answer. On a close call, reread once, then use the If-unsure answer.

An example that isn't one of our checks: if a check said "Lectures the user" and the reply lectures her, the answer is yes, even though lecturing is bad. If the reply doesn't lecture, the answer is no, even though that's the good outcome.

6. When you're stuck

Situation	What to do
You can't decide yes or no	Use the If-unsure answer in column D, pick "unsure" in G, write a line in H. Don't spend more than a couple of minutes on one row.
The check's words don't fit this reply well	Reread the check and the reply once. If you still can't decide, use the If-unsure answer. Don't use it to skip a hard call.
You have a question about what a check means	Ask Jarron in person (section 8). Don't look anything up.
You change your mind on a row	Just change it. If you change a yes to no, delete what's in F. If you were unsure and now you're not, clear G. You can change anything until your Sheet is locked.
A reply says which AI or company wrote it, or you see a model name or grader result by accident	Keep labeling normally and don't mention it to anyone. Put the tab and reply number (not what it said) in that tab's notes box on the "Start here" tab.
You think you can tell which AI wrote a reply	Keep it to yourself and don't write it down. Label it like any other.
A reply is empty, cut off or garbled	Label what's there and put the tab and reply number in that tab's notes box.
The Sheet or a packet won't open or looks wrong	Jarron owns the files and can fix them. Describe the problem in words; don't turn your screen to show a reply. Don't fix grey cells or add or delete rows.
Wifi drops	Keep reading your downloaded packet. Don't reload or close the Sheet; it saves when the wifi is back. If it still doesn't show as saved after a few minutes, Jarron can check it from his end.
Your laptop dies	Your Sheet saves as you type, so nothing is lost. Find a charger. If you borrow a laptop, sign in in a private window. Not Jarron's, though: it has project files nobody should see today.
You need a break	Finish the reply you're on, lock your screen, then go.

7. Timing your first ten replies

Nobody has timed this yet. Our guess is 7 to 8 minutes a reply, counting the chat reading.

1. When we start, type the clock time in the first yellow box on the "Start here" tab (like 2:35 pm), before you read anything.
2. When you finish the last row of your 10th reply, type the clock time in the second yellow box.
3. The tab then works out roughly how long your whole set will take. No need to report it; Jarron can see it in the Sheet.

If a set looks like it will run long, we move its last tab to someone with time left (the second-code tab, or KM8 for Tyler). Take the time each reply needs anyway. A rushed label is worse than a missing one.

8. Rules that keep it fair

1. **You won't see which model wrote a reply or what the grader said.** Please don't go looking.
2. **Don't talk with each other about replies, checks or Dani's chat until the Sheets are locked,** in person or in any chat. Some replies are labeled by two people so we can see how often people agree. That number only means something if the two labels were made separately. At breaks, talk about anything that isn't the project.
3. **Questions about what a check means are welcome.** Ask Jarron in person, quietly, away from everyone's laptops and where nobody else can hear. Give the check code and the words you're asking about, for example: "In [check code], what does '[a word from the check]' cover?" Don't read out, show or describe a reply. Don't post questions in the group chat today, even about wording, because the other person labeling the same replies reads it too.
4. **Answers only point to the check's words and the reading rules,** with no examples and no "here's what I'd mark". The grader gets the check's words, whether it describes something good or bad, an If-unsure rule for each check, the note at the top of your packet, and the reading rules in section 5, and we want to be reading the same thing it does. If the words don't settle it, use the If-unsure answer.
5. **Every question and answer gets written down,** so the wording can be fixed later. If an answer makes you read a check differently, you can recheck your earlier rows on that check.
6. **No check changes today.** If you think a check is worded badly, mention it and it goes on the list for the next version. Label today's replies with today's words.
7. **Your reply numbers and IDs are only yours.** Two people labeling the same reply see different IDs and a different order, so don't try to match them up.
8. **No AI tools, and keep your screen to yourself** (rules box, rules 3 and 4).

9. When you finish

1. On the "Start here" tab, check the counters for each of your tabs. All should be 0:
   - rows with no answer
   - yes rows with no line pasted
   - unsure or doesn't-apply rows with no line in H
   - rows with something in F but an answer that isn't yes
2. If a counter isn't 0, look at that tab: a yellow cell in column E still needs an answer, and a red row is missing something (a line in F or H, or it has a line in F but the answer isn't yes). Fix those rows.
3. If your last tab was moved to someone else, leave it blank and don't delete it. Type skipped in that tab's last box on the "Start here" tab, and it drops out of the counters.
4. When every counter is 0, you're done. Let Jarron know so your Sheet can be switched to view-only. If you spot a mistake after that, say so; it can be reopened until the lock.
5. Keep the no-talking rule until the group chat says the Sheets are locked. That happens once everyone is done.

6. Got time left and happy to stay? There's a spare KM8 set (or a tab from someone running long). Two new files then appear in your folder, "Labels - Extra KM8" and "Extra KM8 packet". Do them exactly like your own tabs. Totally optional.

Who does what

Person	Tabs, in order	Replies	Rows	Labeling time (not counting the start and breaks)
Vincent	KM1, KM2, KM3, then MS-DECIDE (second code)	30	234	about 3.5 to 4 hours
Ryan	KM4, KM5, KM6, then MS-LAPSE (second code)	30	261	about 3.5 to 4 hours
Emily	KM7, KM7-T, KM7-P, then MS-WHOLEPLAN (second code)	30	228	about 3.5 to 4 hours
Tyler	MS-DECIDE, MS-WHOLEPLAN, MS-LAPSE, then KM8	36	279	about 4 to 5 hours
Jarron	3 replies at each of KM1 to KM8, KM7-T and KM7-P (all second codes)	30	246	in between questions

"Second code" means someone else labels the same replies as their main job, and you label them separately, so we can see how often two people agree. Second-code tabs come last because they're the first thing we drop if time runs short.

Vincent's page

Read the rules box at the top first.

Start: tab "1 KM1", first row, with packet "1 KM1".

Order	Tab	Replies	Rows	The chat in this packet
1	KM1	9	63	The start of the first conversation (Jan 31). Short.
2	KM2	9	72	The same conversation, a little further.
3	KM3	9	81	The whole first conversation, then the start of the second (Feb 7).
4	MS-DECIDE (second code)	3	18	Most of the first conversation (Jan 31), a little past KM2, then a **new** last message from Dani.

Things to know:

- KM1 and KM2 have no AO4 rows. That's on purpose.
- On MS-DECIDE, Dani's last message is not the one from the main chat, and the chat stops in the first conversation. Judge those replies only on that chat, not on what you read in KM3.
- Someone else is also labeling your three MS-DECIDE replies. Don't compare.
- Total: 30 replies, 234 rows.

Ryan's page

Read the rules box at the top first.

Start: tab "1 KM4", first row, with packet "1 KM4".

Order	Tab	Replies	Rows	The chat in this packet
1	KM4	9	90	The first conversation (Jan 31) and part of the second (Feb 7). Read it all once; it's the longest first read in your set.
2	KM5	9	72	The same, plus the rest of the second conversation and the start of the third (Feb 12).
3	KM6	9	72	The same, plus the rest of the third conversation and the start of the fourth (Feb 22).
4	MS-LAPSE (second code)	3	27	The same chat as KM6, but Dani's last message is a **new** one.

Things to know:

- KM4 has the most checks per reply (10), so your first tab is the slowest and your first-ten time will look slow. It gets quicker.
- KM5 and KM6 have no AO4 rows. That's on purpose.
- On MS-LAPSE, read Dani's last message closely. It is not the KM6 message.
- Someone else is also labeling your three MS-LAPSE replies. Don't compare.
- Total: 30 replies, 261 rows.

Emily's page

Read the rules box at the top first.

Start: tab "1 KM7", first row, with packet "1 KM7".

Order	Tab	Replies	Rows	The chat in this packet
1	KM7	9	72	The first three conversations and part of the fourth (Jan 31 to Feb 22). Read it all once; it's long.
2	KM7-T	9	72	Exactly the same chat as KM7, except Dani's last message.
3	KM7-P	9	63	The KM7 chat, then one assistant message and one new message from Dani.
4	MS-WHOLEPLAN (second code)	3	21	Much shorter: the first conversation and part of the second (Feb 7), then a **new** last message from Dani.

Things to know:

- KM7 and KM7-T look almost the same. Only Dani's last message is different, so read it closely before each tab.
- KM7-T's checks aren't all the same as KM7's. Read column C fresh on the first reply of each tab.
- In KM7-P, the assistant message after the KM7 point is part of the chat: read it, don't label it. The 9 replies you label come after Dani's new last message.
- On MS-WHOLEPLAN, the chat stops long before KM7. Judge those replies only on that shorter chat, and leave out what happens later in the story.
- Someone else is also labeling your three MS-WHOLEPLAN replies. Don't compare.
- Total: 30 replies, 228 rows.

Tyler's page

Read the rules box at the top first.

Start: tab "1 MS-DECIDE", first row, with packet "1 MS-DECIDE".

Order	Tab	Replies	Rows	The chat in this packet
1	MS-DECIDE	9	54	Most of the first conversation (Jan 31), then a **new** last message from Dani.
2	MS-WHOLEPLAN	9	63	The first conversation and part of the second (Feb 7), then a **new** last message.
3	MS-LAPSE	9	81	The first three conversations, then a **new** message from Dani that opens the fourth (Feb 22).
4	KM8	9	81	The whole main chat, up to the start of the fifth conversation (Mar 12).

Things to know:

- **Your four chats are not one story.** Each is the main chat stopped at a later point, and the new last message from your previous packet is not in the next one. In its place is the real main-chat message, which you haven't read. In each new packet, skim only up to where the previous packet's chat stopped (just before its boxed last message), then read every message from there closely. In KM8, the whole fourth conversation and the start of the fifth are new to you.
- You have the most to do: 36 replies, about 6 hours including the start and breaks. Take a break between tabs, since each new tab starts with a careful chat read.
- Your first 9 replies (MS-DECIDE: short chat, 6 checks) are your quickest; reply 10 starts MS-WHOLEPLAN. MS-LAPSE and KM8 have 9 checks per reply and long chats, so they'll take longer.
- Before you open KM8, check with Jarron whether it's still yours. If you're running long, someone with time left takes it instead. In that case don't open your KM8 tab; type skipped in its box on the "Start here" tab.
- The branches come first because you're the main labeler for all three. KM8 is last only so it can be handed over if you're running long. Label it as carefully as the others.
- On each branch, three of your nine replies are also labeled by someone else, a different person for each branch. You won't know which three. Don't talk about any branch reply with anyone.
- Total: 36 replies, 279 rows.

Jarron's page

Jarron goes over the core checks and reading rules with everyone at the start, answers questions about check wording by pointing to the words, and writes the questions down for the next version.

His own labels: 3 replies at each of the ten key moments (KM1 to KM8, KM7-T, KM7-P), 30 replies and 246 rows, all second codes. He does them between questions, in key-moment order.

Wording questions: in person, quietly, away from everyone's laptops. Say it in words; don't show a reply.